

Exploratory Data Analysis (EDA) is a critical step in data analysis that involves using various techniques to understand the structure, patterns, relationships, and potential issues in a dataset before applying statistical models or machine learning algorithms. The goal of EDA is to explore and summarize the main characteristics of the data, often with visual methods, to gain insights and help guide further analysis.

Here are the key concepts in EDA theory:

## 1. Understanding the Dataset

1.Data Types: EDA begins by understanding the types of variables in the dataset (e.g., numerical, categorical, ordinal, etc.). Each type will require different techniques for analysis.

2.Missing Values: Identifying missing data points and deciding how to handle them (e.g., removing, imputing, or leaving them as-is).

3.Outliers: Detecting data points that significantly differ from others, which may indicate errors, anomalies, or important variations in the data.

4.Data Distribution: Understanding the distribution of the data (e.g., normal, skewed, bimodal) helps in deciding which statistical techniques are appropriate.

## 2. Statistical Summaries

1.Central Tendency: Measures such as the mean, median, and mode are used to understand the "center" of the data.

2.Dispersion: Measures such as variance, standard deviation, and interquartile range (IQR) help assess how spread out the data is.

3.Shape of Distribution: Skewness (asymmetry) and kurtosis (peakedness) help characterize the distribution of the data.

4.Correlations: Identifying relationships between variables using correlation coefficients (e.g., Pearson or Spearman) to understand associations.

## 3. Visualization Techniques

1.Histograms: Used to visualize the distribution of a single variable (especially numerical data).

2.Box Plots: Visualize the spread and detect outliers, displaying the median, quartiles, and potential outliers.

3.Bar Charts: For categorical variables, bar charts show the frequency or proportion of each category.

4. Scatter Plots: Useful for visualizing relationships between two numerical variables.
5. Pair Plots: Help visualize relationships among multiple numerical variables simultaneously.
6. Heatmaps: Used to visualize correlation matrices or missing data patterns.
7. Violin Plots: Combine aspects of box plots and density plots to show the distribution of data.

## 4. Handling Missing Data

1. Deletion: Removing rows or columns with missing data if it's minimal.
2. Imputation: Replacing missing values with estimates, such as the mean, median, mode, or a more sophisticated imputation method (e.g., using algorithms).
3. Marking Missing: Creating a separate indicator column that marks where data is missing.

## 5. Feature Engineering

1. Transformations: Sometimes, raw data needs to be transformed (e.g., scaling, normalization, or log transformations) to make it suitable for modeling.
2. Encoding Categorical Variables: Categorical data might be encoded into numerical forms (e.g., one-hot encoding, label encoding) for use in machine learning algorithms.

## 6. Dimensionality Reduction

.If the dataset has many features, dimensionality reduction techniques (such as Principal Component Analysis, PCA) may be applied to reduce complexity and highlight key features.

## 7. Identifying Patterns and Relationships

1. Clustering: Grouping data points into clusters based on similarity. This helps identify natural groupings within the data.
2. Trend Detection: Looking for trends or patterns over time or across different subgroups (e.g., seasonal patterns in time series data)
3. Anomaly Detection: Identifying data points that deviate significantly from expected patterns (outliers).

## 8. Feature Selection

After performing EDA, you may want to focus on the most relevant features for modeling. This is typically done by considering correlations, feature importance scores, or domain knowledge.

## 9. Data Transformation

1. Scaling and Normalization: Ensuring that numerical variables are on the same scale (especially important for algorithms like k-nearest neighbors or gradient descent).

2. Log Transformations: Used to handle skewed data or make distributions more normal.

## 10. Communication of Findings

The results from EDA should be clearly communicated using graphs, tables, and concise summaries to inform further analysis and decision-making.

## Tools and Libraries for EDA

There are several tools and libraries commonly used for EDA, especially in Python:

Pandas: For data manipulation and basic statistical summaries.

Matplotlib & Seaborn: For creating static visualizations like histograms, scatter plots, and box plots.

Plotly: For interactive visualizations.

Scikit-learn: For more advanced statistical techniques and dimensionality reduction methods.

```
In [ ]: # import libraries
import numpy as np
import pandas as pd

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
In [65]: import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as st
%matplotlib inline

sns.set(style="whitegrid")
```

```
In [67]: import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: # import dataset
```

```
In [69]: df = pd.read_csv(r"C:\Users\navvee\OneDrive\Desktop\heart.csv")
```

```
In [71]: df
```

```
Out[71]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	tl
0	63	1	3	145	233	1	0	150	0	2.3	0	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	
299	45	1	3	110	264	0	1	132	0	1.2	1	0	
300	68	1	0	144	193	1	1	141	0	3.4	1	2	
301	57	1	0	130	131	0	1	115	1	1.2	1	1	
302	57	0	1	130	236	0	0	174	0	0.0	1	1	

303 rows × 14 columns



```
In [77]: #exploratory data analysis
print('The shape of the dataset : ', df.shape)
```

The shape of the dataset : (303, 14)

```
In [ ]:
```

```
In [79]: df.head()
```

```
Out[79]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2



```
In [81]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null   int64
 1   sex         303 non-null   int64
 2   cp          303 non-null   int64
 3   trestbps    303 non-null   int64
 4   chol        303 non-null   int64
 5   fbs         303 non-null   int64
 6   restecg     303 non-null   int64
 7   thalach     303 non-null   int64
 8   exang       303 non-null   int64
 9   oldpeak     303 non-null   float64
10   slope       303 non-null   int64
11   ca          303 non-null   int64
12   thal        303 non-null   int64
13   target      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

In [83]: `df.dtypes`

```

Out[83]: age         int64
sex         int64
cp          int64
trestbps    int64
chol        int64
fbs         int64
restecg     int64
thalach     int64
exang       int64
oldpeak     float64
slope       int64
ca          int64
thal        int64
target      int64
dtype: object

```

In [85]: `df.describe()`

```

Out[85]:
```

	age	sex	cp	trestbps	chol	fbs	restecg
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528000
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525000
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000
<b>50%</b>	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000
<b>75%</b>	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000

```
In [87]: df.columns
```

```
Out[87]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',  
              'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],  
              dtype='object')
```

```
In [89]: df['target'].nunique()
```

```
Out[89]: 2
```


```
In [91]: df['target'].nunique()
```

```
Out[91]: 2
```

```
In [93]: df['target'].value_counts()
```

```
Out[93]: target  
1      165  
0      138  
Name: count, dtype: int64
```

```
In [125... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(x="target", data=df)  
plt.show()
```




```
In [99]: df.groupby('sex')['target'].value_counts()
```


```
Out[99]: sex  target
0      1      72
        0      24
1      0     114
        1      93
Name: count, dtype: int64
```

```
In [103... f, ax = plt.subplots(figsize=(8, 6))
ax = sns.countplot(x="sex", hue="target", data=df)
plt.show()
```


<Figure size 640x480 with 0 Axes>



```
In [127... ax = sns.catplot(x="target", col="sex", data=df, kind="count", height=5, aspect=
plt.show()
```




```
In [129... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(y="target", hue="sex", data=df)  
plt.show()
```




```
In [131... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(x="target", data=df, palette="Set3")  
plt.show()
```






```
In [133... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(x="target", data=df, facecolor=(0, 0, 0, 0), linewidth=5, edge  
plt.show()
```




```
In [135... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(x="target", hue="fbs", data=df)
```

```
plt.show()
```



```
In [137... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(x="target", hue="exang", data=df)  
plt.show()
```



```
In [139... correlation = df.corr()
```

```
In [141... correlation['target'].sort_values(ascending=False)
```

```
Out[141... target      1.000000  
cp          0.433798  
thalach     0.421741  
slope       0.345877  
restecg     0.137230  
fbs         -0.028046  
chol        -0.085239  
trestbps    -0.144931  
age         -0.225439  
sex         -0.280937  
thal        -0.344029  
ca          -0.391724  
oldpeak     -0.430696  
exang       -0.436757  
Name: target, dtype: float64
```


```
In [143... df['cp'].nunique()
```

```
Out[143... 4
```

```
In [145... df['cp'].value_counts()
```

```
Out[145... cp  
0      143  
2       87  
1       50  
3       23  
Name: count, dtype: int64
```


```
In [147... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.countplot(x="cp", data=df)  
plt.show()
```



```
In [149... df.groupby('cp')['target'].value_counts()
```

```
Out[149... cp target
0  0      104
   1       39
1  1       41
   0        9
2  1       69
   0       18
3  1       16
   0        7
Name: count, dtype: int64
```

```
In [151... f, ax = plt.subplots(figsize=(8, 6))
ax = sns.countplot(x="cp", hue="target", data=df)
plt.show()
```



In [159...

```
ax = sns.catplot(x="target", col="cp", data=df, kind="count", height=8, aspect=1
plt.show()
```



In [161...


```
df['thalach'].nunique()
```

Out[161...


91

In [163...


```
f, ax = plt.subplots(figsize=(10,6))
x = df['thalach']
ax = sns.distplot(x, bins=10)
plt.show()
```




```
In [165... f, ax = plt.subplots(figsize=(10,6))
x = df['thalach']
x = pd.Series(x, name="thalach variable")
ax = sns.distplot(x, bins=10)
plt.show()
```




```
In [167... f, ax = plt.subplots(figsize=(10,6))
x = df['thalach']
ax = sns.distplot(x, bins=10, vertical=True)
plt.show()
```




```
In [169... f, ax = plt.subplots(figsize=(10,6))
x = df['thalach']
x = pd.Series(x, name="thalach variable")
ax = sns.kdeplot(x)
plt.show()
```



```
In [171... f, ax = plt.subplots(figsize=(10,6))
x = df['thalach']
x = pd.Series(x, name="thalach variable")
ax = sns.kdeplot(x, shade=True, color='r')
plt.show()
```




```
In [173... f, ax = plt.subplots(figsize=(10,6))
x = df['thalach']
ax = sns.distplot(x, kde=False, rug=True, bins=10)
plt.show()
```




```
In [175... f, ax = plt.subplots(figsize=(8, 6))
sns.stripplot(x="target", y="thalach", data=df)
plt.show()
```






```
In [177... f, ax = plt.subplots(figsize=(8, 6))  
sns.stripplot(x="target", y="thalach", data=df, jitter = 0.01)  
plt.show()
```




```
In [179... f, ax = plt.subplots(figsize=(8, 6))  
sns.boxplot(x="target", y="thalach", data=df)
```

```
plt.show()
```




```
In [181... plt.figure(figsize=(16,12))
plt.title('Correlation Heatmap of Heart Disease Dataset')
a = sns.heatmap(correlation, square=True, annot=True, fmt='.2f', linecolor='white')
a.set_xticklabels(a.get_xticklabels(), rotation=90)
a.set_yticklabels(a.get_yticklabels(), rotation=30)
plt.show()
```



In [184...

```
num_var = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak', 'target' ]
sns.pairplot(df[num_var], kind='scatter', diag_kind='hist')
plt.show()
```




```
In [186... df['age'].unique()
```

```
Out[186... 41
```


```
In [188... df['age'].describe()
```

```
Out[188... count    303.000000
mean      54.366337
std       9.082101
min       29.000000
25%      47.500000
50%      55.000000
75%      61.000000
max       77.000000
Name: age, dtype: float64
```


```
In [190... f, ax = plt.subplots(figsize=(10,6))
x = df['age']
ax = sns.distplot(x, bins=10)
plt.show()
```




```
In [192... f, ax = plt.subplots(figsize=(8, 6))
sns.stripplot(x="target", y="age", data=df)
plt.show()
```




```
In [194... f, ax = plt.subplots(figsize=(8, 6))
sns.boxplot(x="target", y="age", data=df)
plt.show()
```




```
In [196... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.scatterplot(x="age", y="trestbps", data=df)  
plt.show()
```




```
In [198... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.regplot(x="age", y="trestbps", data=df)  
plt.show()
```



```
In [200... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.scatterplot(x="age", y="chol", data=df)  
plt.show()
```




```
In [202... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.regplot(x="age", y="chol", data=df)  
plt.show()
```




```
In [204... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.scatterplot(x="chol", y = "thalach", data=df)
```



```
plt.show()
```



```
In [206... f, ax = plt.subplots(figsize=(8, 6))  
ax = sns.regplot(x="chol", y="thalach", data=df)  
plt.show()
```



```
In [208... df.isnull().sum()
```

```
Out[208... age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```


```
In [210... assert pd.notnull(df).all().all()
```

```
In [212... assert (df >= 0).all().all()
```

```
In [214... df['age'].describe()
```

```
Out[214... count    303.000000
mean      54.366337
std       9.082101
min       29.000000
25%      47.500000
50%      55.000000
75%      61.000000
max       77.000000
Name: age, dtype: float64
```


```
In [216... f, ax = plt.subplots(figsize=(8, 6))
sns.boxplot(x=df["age"])
plt.show()
```



```
In [218...] df['trestbps'].describe()
```

```
Out[218...] count    303.000000
mean      131.623762
std       17.538143
min       94.000000
25%      120.000000
50%      130.000000
75%      140.000000
max       200.000000
Name: trestbps, dtype: float64
```


```
In [220...] f, ax = plt.subplots(figsize=(8, 6))
sns.boxplot(x=df["trestbps"])
plt.show()
```



```
In [222...] df['chol'].describe()
```

```
Out[222...] count    303.000000  
mean      246.264026  
std        51.830751  
min       126.000000  
25%       211.000000  
50%       240.000000  
75%       274.500000  
max       564.000000  
Name: chol, dtype: float64
```


```
In [224...] f, ax = plt.subplots(figsize=(8, 6))  
sns.boxplot(x=df["chol"])  
plt.show()
```



```
In [226...] df['thalach'].describe()
```

```
Out[226...] count    303.000000
mean      149.646865
std       22.905161
min       71.000000
25%      133.500000
50%      153.000000
75%      166.000000
max       202.000000
Name: thalach, dtype: float64
```


```
In [228...] f, ax = plt.subplots(figsize=(8, 6))
sns.boxplot(x=df["thalach"])
plt.show()
```



```
In [230...] df['oldpeak'].describe()
```

```
Out[230...] count    303.000000
mean         1.039604
std          1.161075
min          0.000000
25%          0.000000
50%          0.800000
75%          1.600000
max          6.200000
Name: oldpeak, dtype: float64
```

```
In [232...] f, ax = plt.subplots(figsize=(8, 6))
sns.boxplot(x=df["oldpeak"])
plt.show()
```



## Why EDA is Important:

\*Identify Data Quality Issues: Helps in detecting errors, missing values, outliers, or inconsistencies in the dataset.

\*Guide Further Analysis: By exploring the data first, you can make informed decisions about which statistical tests or machine learning algorithms to apply.

\*Visual Insight: Visualization allows you to identify patterns, trends, and relationships that might not be obvious in raw data.

\*Prevent Overfitting: It helps you understand the characteristics of the dataset, preventing you from overfitting your models.

EDA is not just about summarizing data but using those summaries to gain deeper insights, test assumptions, and prepare the data for further analysis or modeling.

## Summary

In summary, EDA is both an art and a science. It involves iteratively exploring and visualizing data to better understand its structure and patterns, which helps in making informed decisions about further analysis, feature engineering, and modeling. The insights gained during EDA are foundational for any data science or statistical modeling task.

In [ ]: